

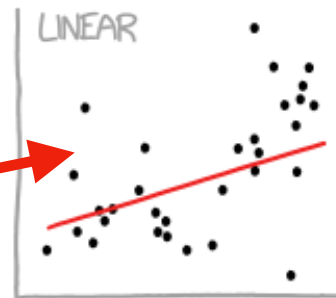
**Welcome to Week #14!**

# K-Nearest Neighbors

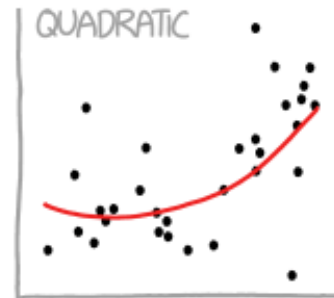
# **K-Nearest Neighbors**

**First: an intro to overfitting**

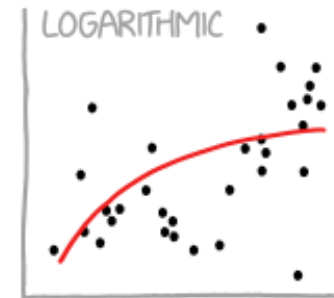
# CURVE-FITTING METHODS AND THE MESSAGES THEY SEND



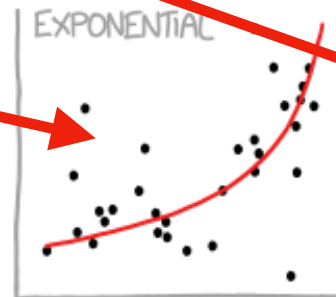
"HEY, I DID A REGRESSION."



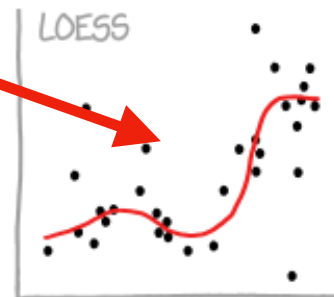
"I WANTED A CURVED LINE, SO I MADE ONE WITH MATH."



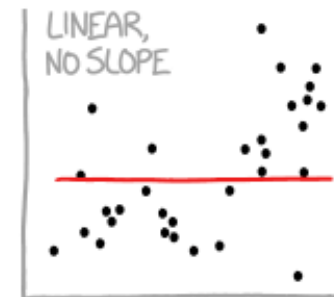
"LOOK, IT'S TAPERING OFF!"



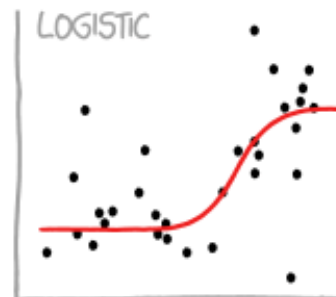
"LOOK, IT'S GROWING UNCONTROLLABLY!"



"I'M SOPHISTICATED, NOT LIKE THOSE BUMBLING POLYNOMIAL PEOPLE."



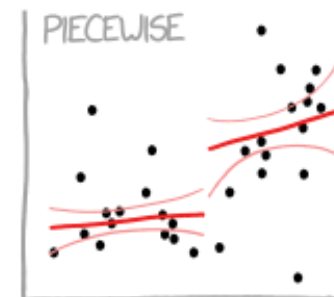
"I'M MAKING A SCATTER PLOT BUT I DON'T WANT TO."



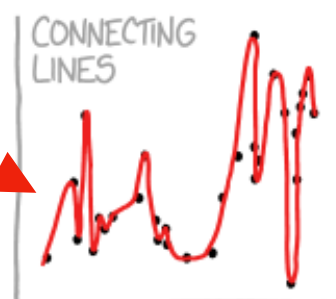
"I NEED TO CONNECT THESE TWO LINES, BUT MY FIRST IDEA DIDN'T HAVE ENOUGH MATH."



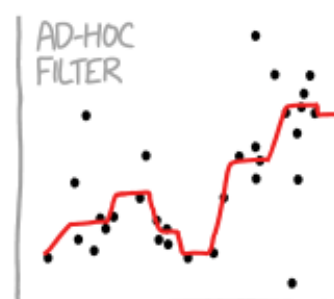
"LISTEN, SCIENCE IS HARD. BUT I'M A SERIOUS PERSON DOING MY BEST."



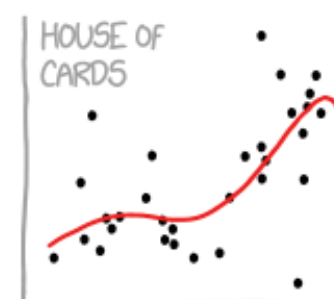
"I HAVE A THEORY, AND THIS IS THE ONLY DATA I COULD FIND."



"I CLICKED 'SMOOTH LINES' IN EXCEL."



"I HAD AN IDEA FOR HOW TO CLEAN UP THE DATA. WHAT DO YOU THINK?"



"AS YOU CAN SEE, THIS MODEL SMOOTHLY FITS THE— WAIT NO NO DON'T EXTEND IT AAAAAA!!!"

Overfit?

Overfit?

Overfit

Underfit?

Is overfitting or underfitting worse?

# Bias-Variance Trade-Off (First Glance)

$$E \left( y_0 - \hat{f}(x_0) \right)^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon)$$

# Bias-Variance Trade-Off (First Glance)

$$E \left( y_0 - \hat{f}(x_0) \right)^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon)$$

**mean square error** if we kept estimating response variable  $y$  by our fitted function of our explanatory variables,  $f(x)$  with different sample datasets at point  $x_0$

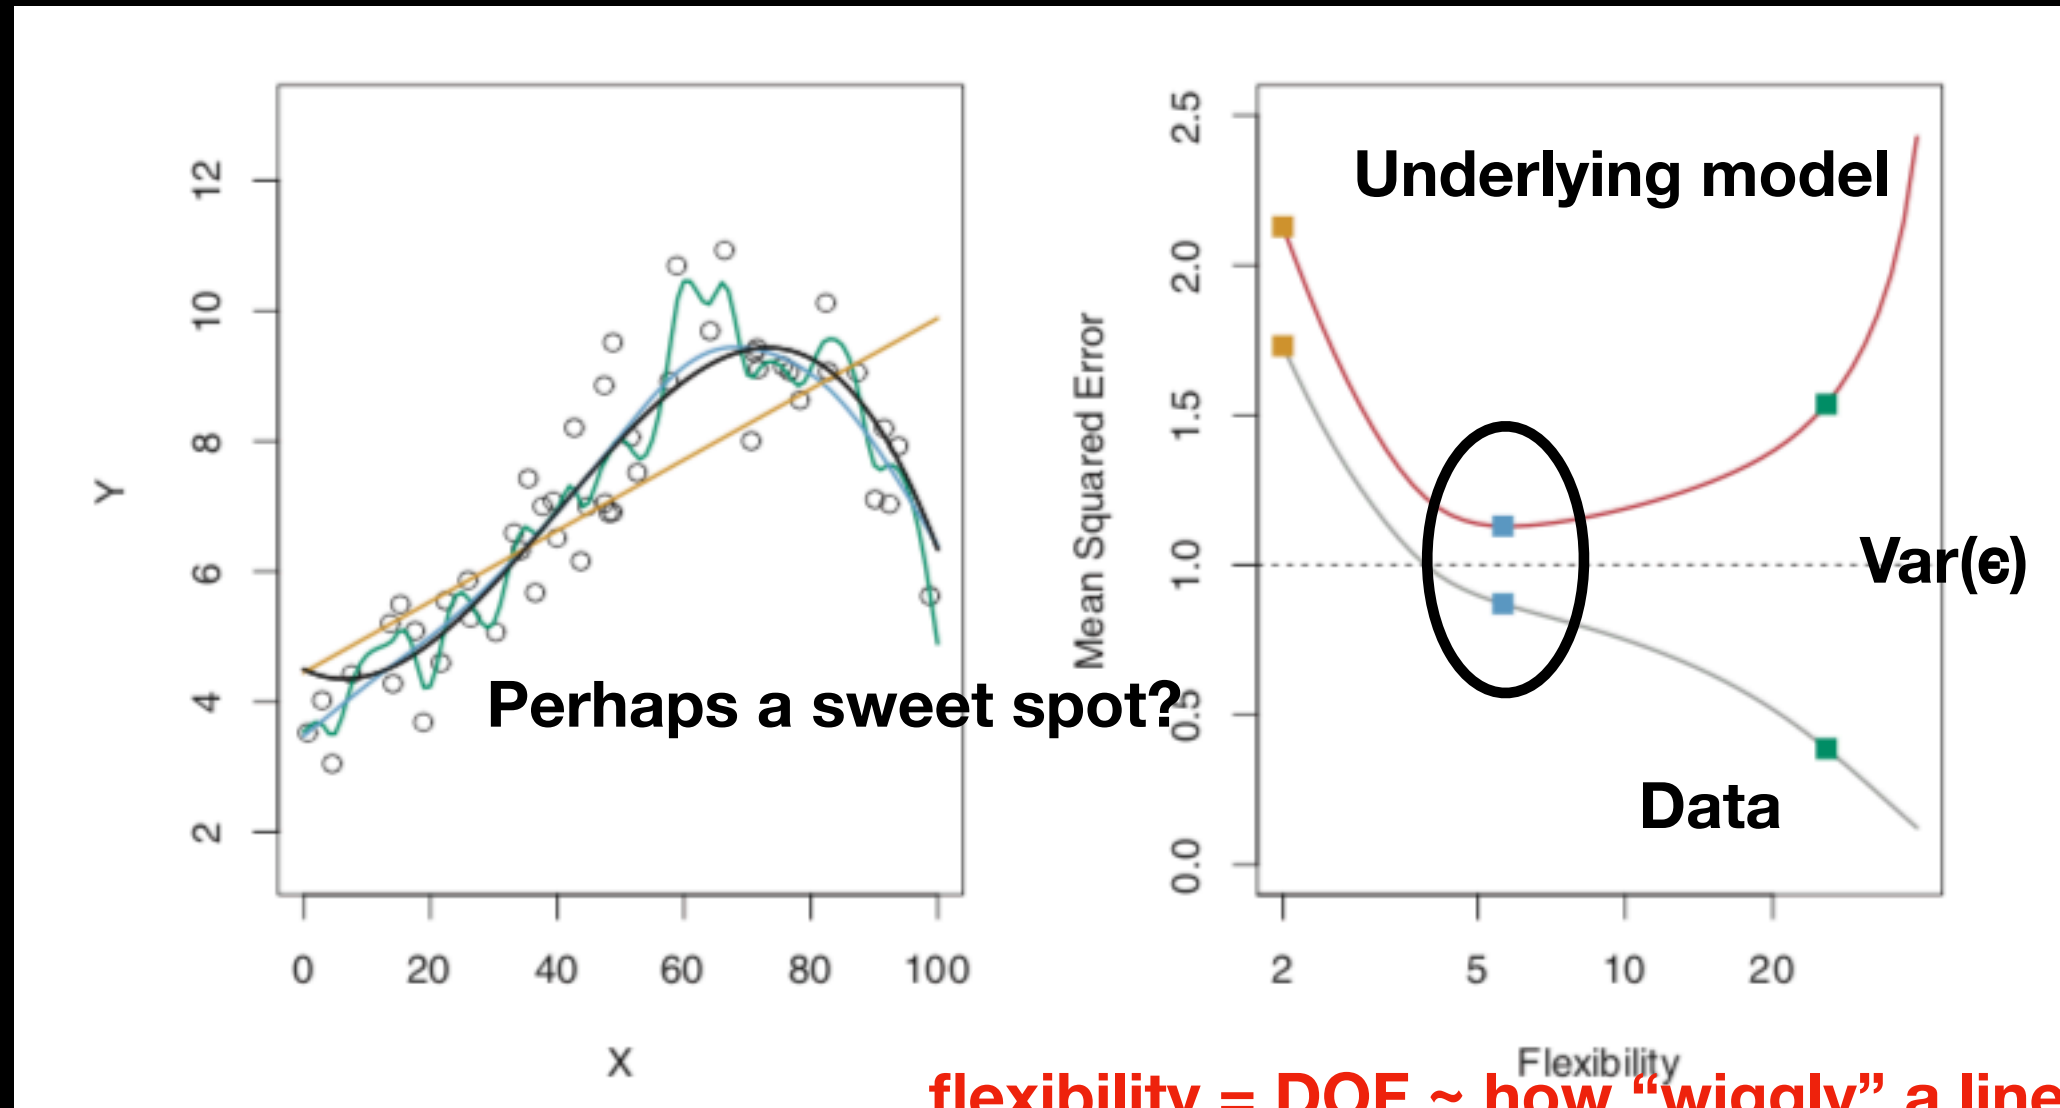
how much our function,  $f$ , changes if we use a different random sample  
(**variance**)

Inherent error (**bias**) in the fact that any model is only an approximation to reality

Inherent error in our measurements

# Bias-Variance Trade-Off (First Glance)

$$E \left( y_0 - \hat{f}(x_0) \right)^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon).$$

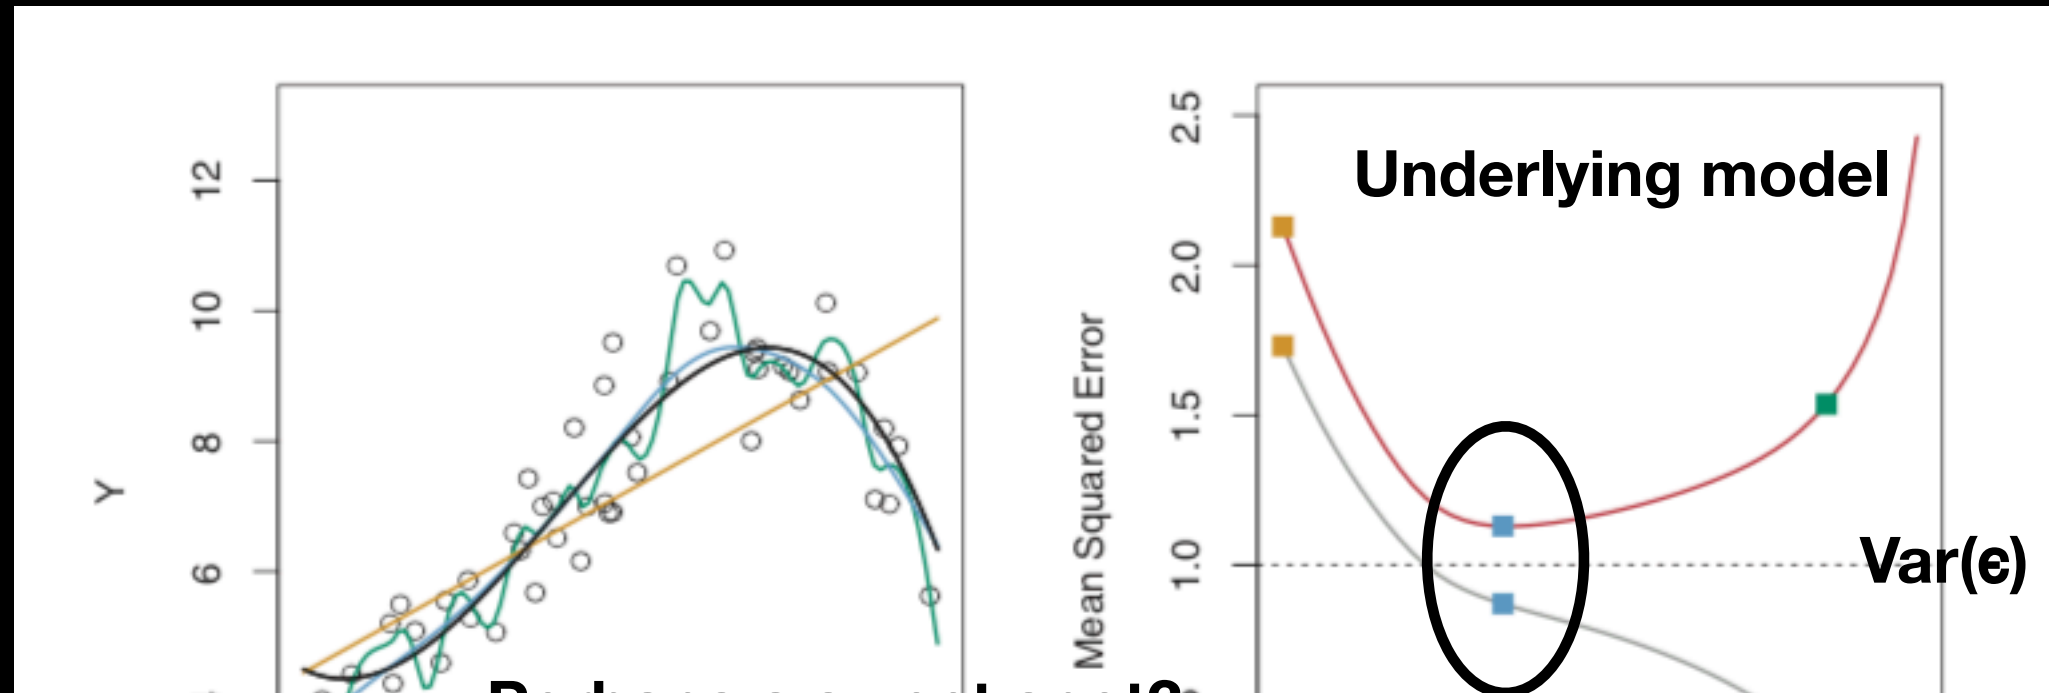


- Actual underlying function -  $y$
- o Simulated data with added error ( $\epsilon$ )
- Linear fit
- Low “flexibility” smooth spline
- High “flexibility” smooth spline

fits data well, but underlying model badly

# Bias-Variance Trade-Off (First Glance)

$$E \left( y_0 - \hat{f}(x_0) \right)^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon).$$



Keep this idea of under/over fitting in mind as we move forward...

- Actual underlying function -  $y$
- o Simulated data with added error ( $\epsilon$ )
- Linear fit
- Low “flexibility” smooth spline
- High “flexibility” smooth spline

flexibility = DOF ~ how “wiggly” a line is

fits data well, but underlying model badly



# K-Nearest Neighbors

~~First: an intro to overfitting~~

# So far...

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_n x_n$$

*predicted y*

*intercept*

# So far...

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_n x_n$$

*predicted y*

*intercept*

So far we've been saying:  
"I have a basic idea of what the functional form of y looks like (a line or a logit) - computer go find the parameters of that functional form."

# So far...

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_n x_n$$

*predicted y*

*intercept*

So far we've been saying:

"I have a basic idea of what the functional form of y looks like (a line or a logit) - computer go find the parameters of that functional form."

This is nice because we have some hope of gaining intuition from our models.

# So far...

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_n x_n$$

*predicted y*

*intercept*

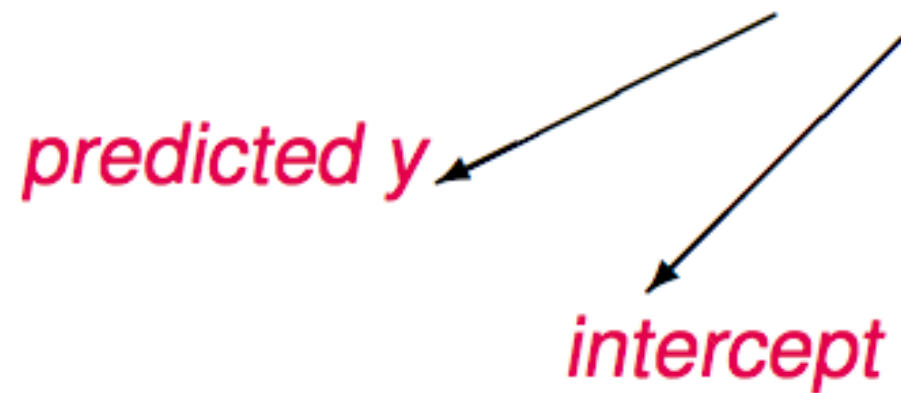
So far we've been saying:  
"I have a basic idea of what the functional form of y looks like (a line or a logit) - computer go find the parameters of that functional form."

This is nice because we have some hope of gaining intuition from our models.

# Now we classify...

# So far...

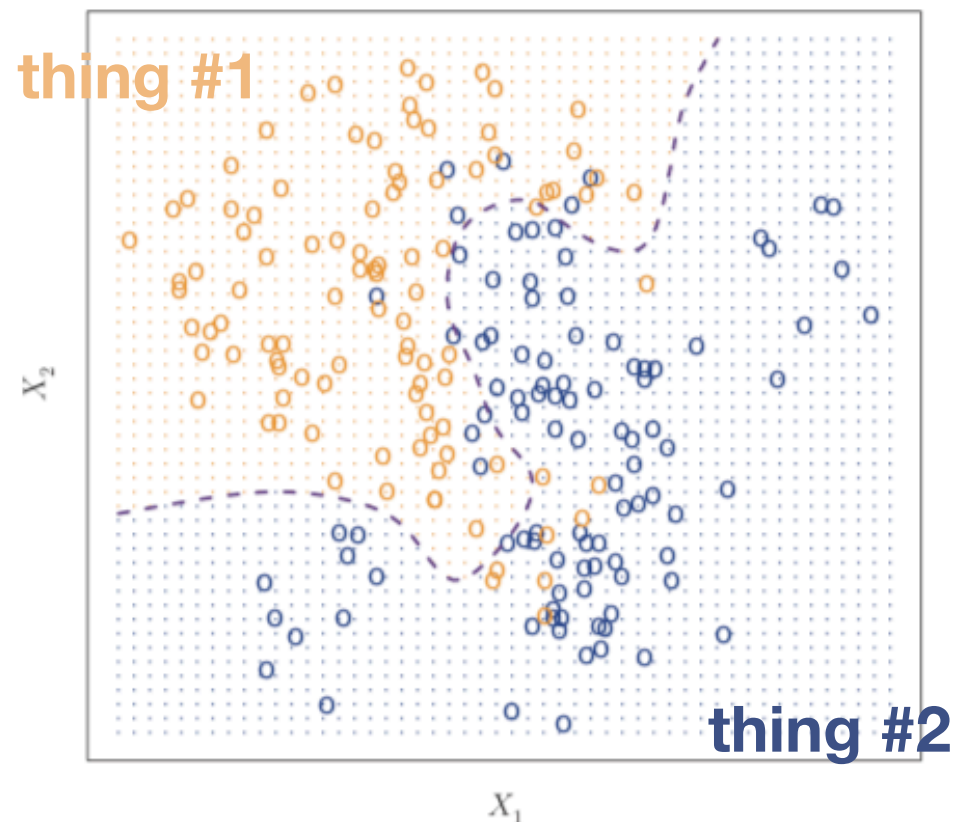
$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_n x_n$$



So far we've been saying:  
“I have a basic idea of what the functional form of  $y$  looks like (a line or a logit) - computer go find the parameters of that functional form.”

This is nice because we have some hope of gaining intuition from our models.

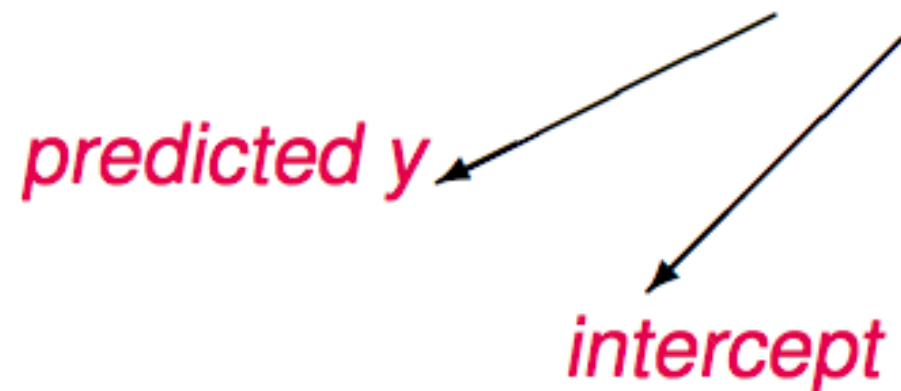
# Now we classify...



“I want to know where thing #1 and thing #2 live in some 2D space - computer, go figure out the boundary between these two things and let me know”

# So far...

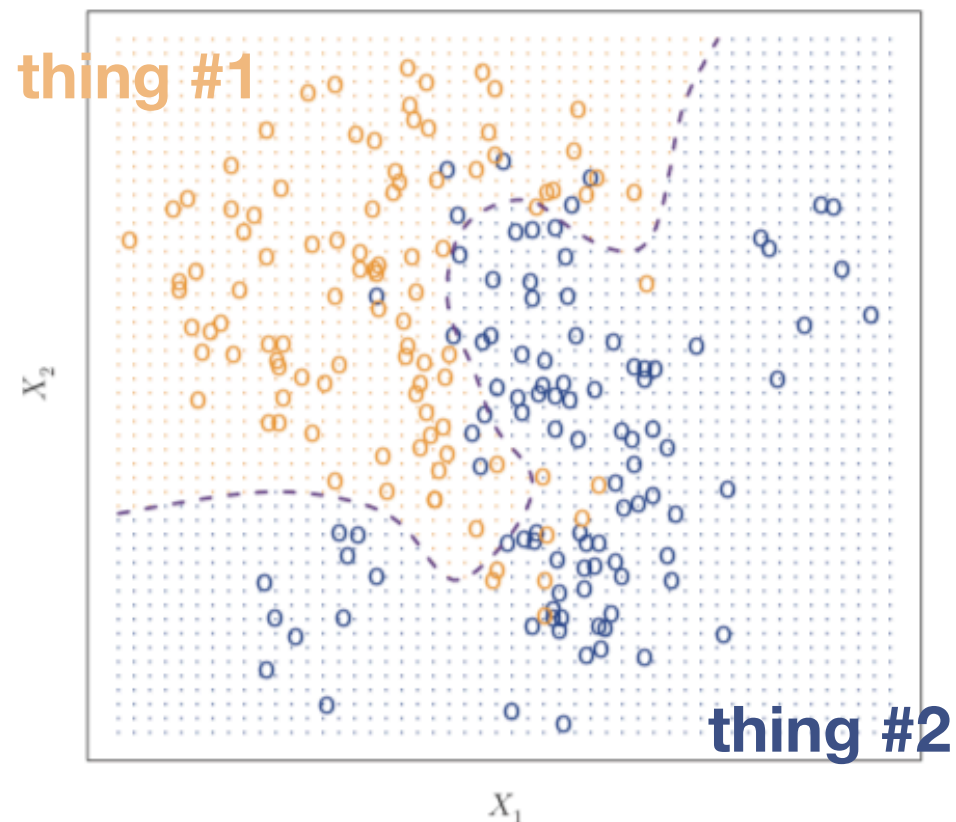
$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_n x_n$$



So far we've been saying:  
"I have a basic idea of what the functional form of y looks like (a line or a logit) - computer go find the parameters of that functional form."

This is nice because we have some hope of gaining intuition from our models.

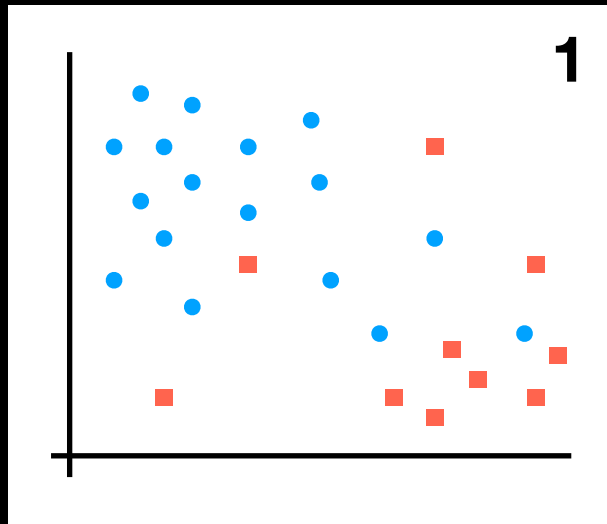
# Now we classify...



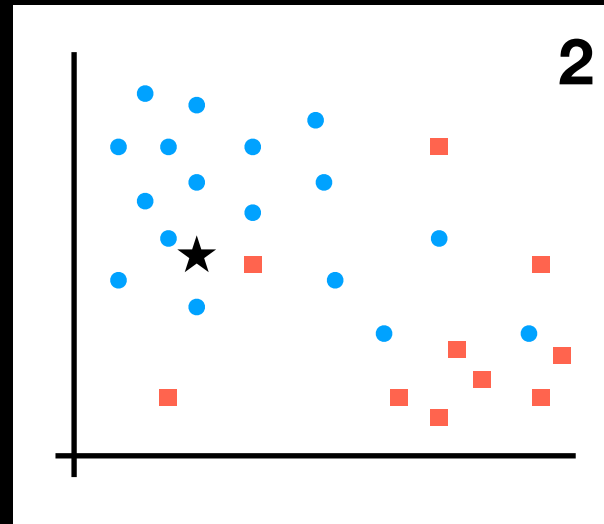
"I want to know where thing #1 and thing #2 live in some 2D space - computer, go figure out the boundary between these two things and let me know"

This is nice because we don't have to assume some model beforehand.

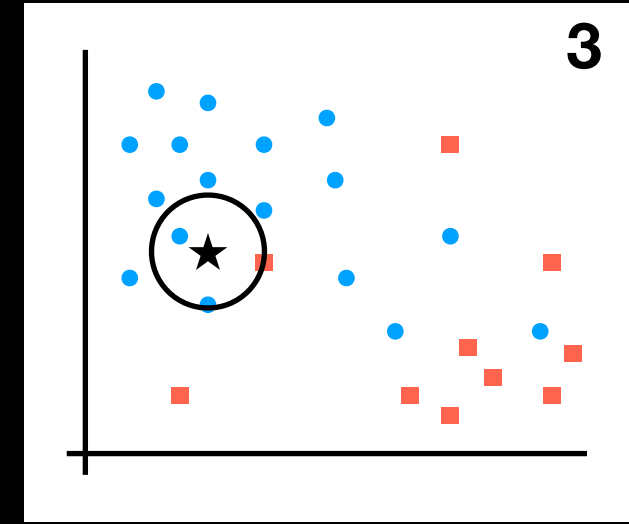
# K Nearest Neighbors, in pictures



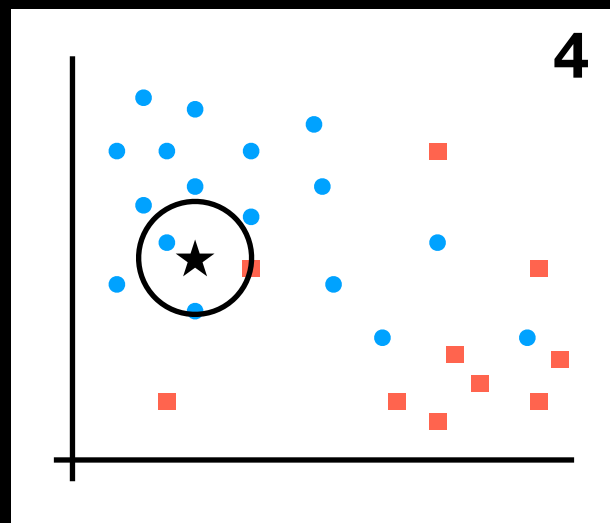
Sample (training) data  
representing underlying  
population



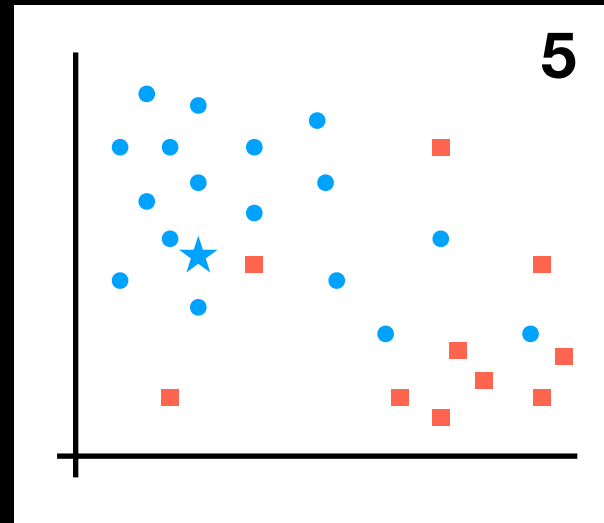
New point of interest



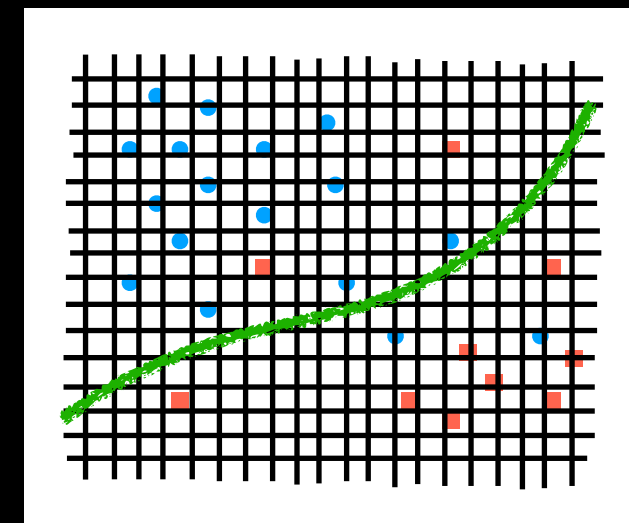
Find k nearest  
neighbors  
(here  $k = 3$ )



count "types" - here  
2/3 points are blue  
 $P(\text{blue}) = 2/3$   
 $P(\text{red}) = 1/3$



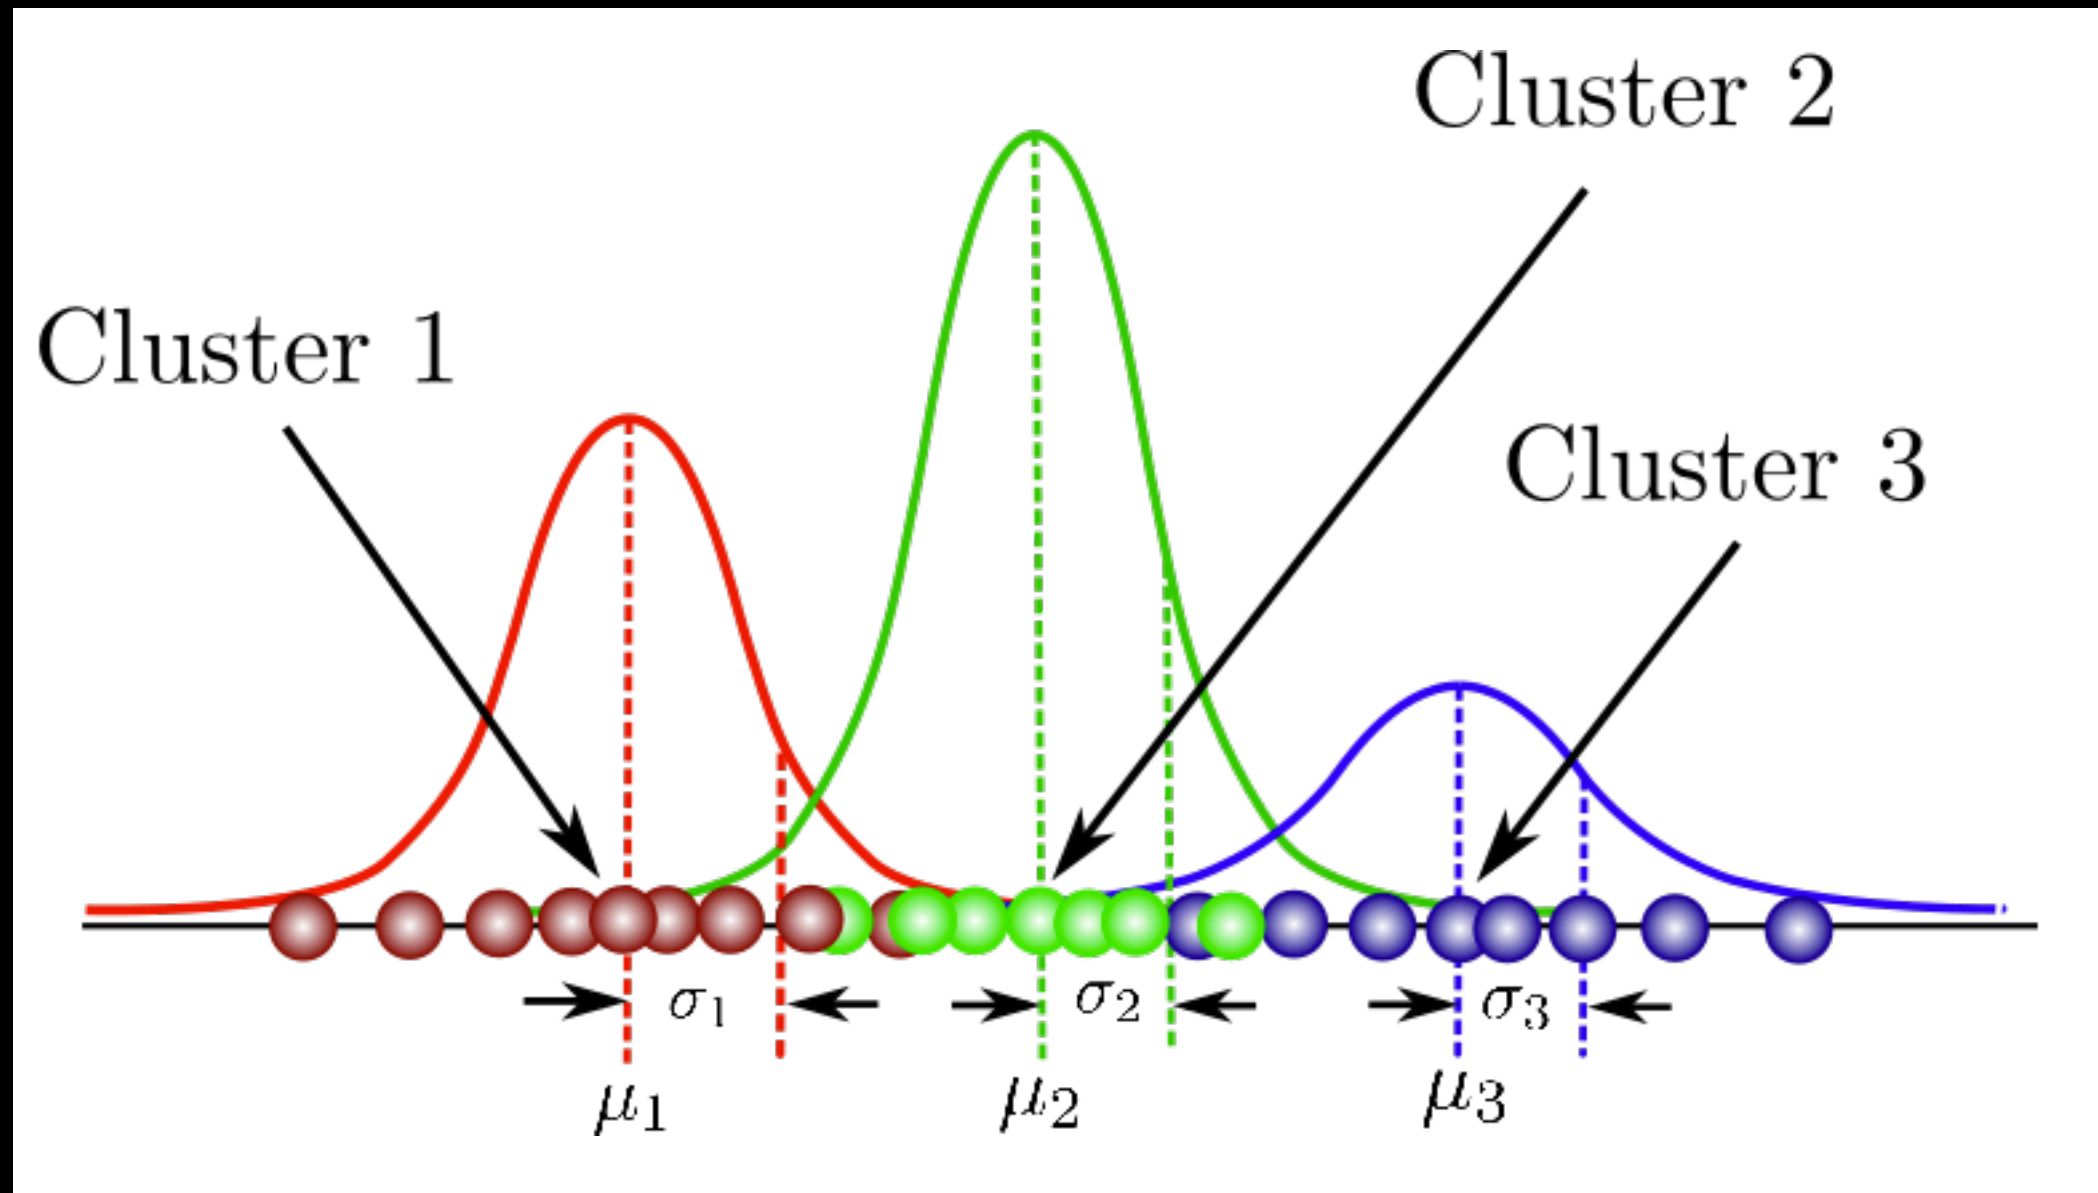
if  $P > \text{cut off}$  say new  
point is in that group  
here:  $P(\text{blue}) > 0.5$



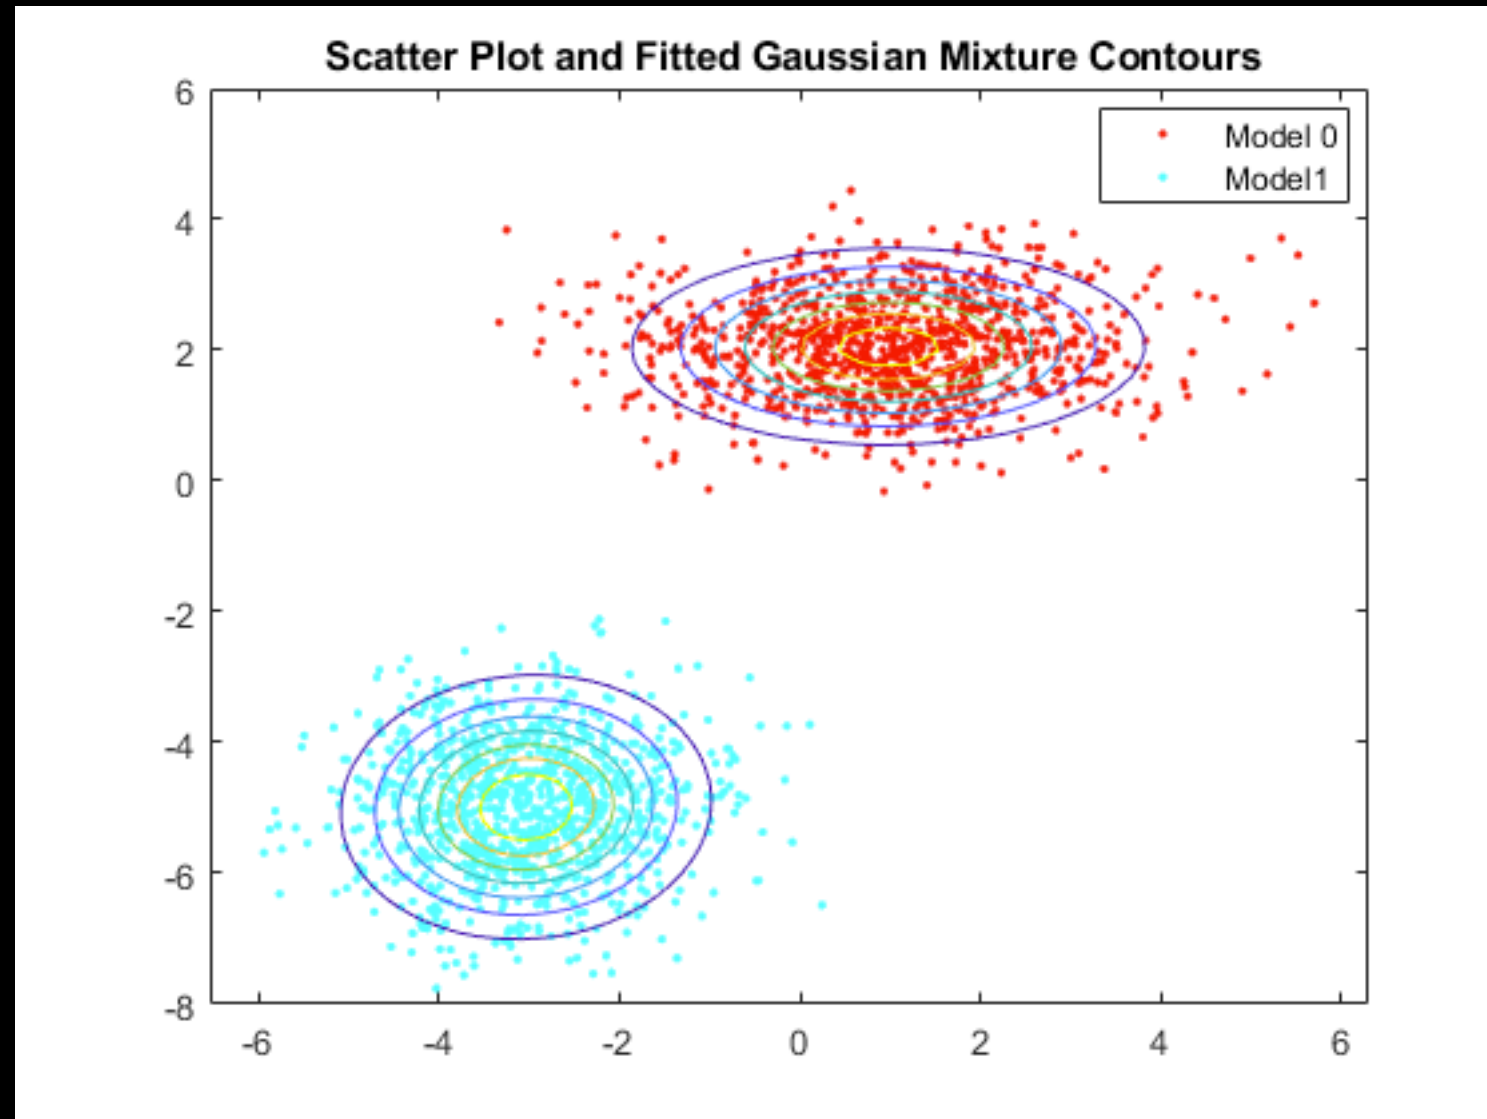
Repeat 2-3 on a grid  
& draw a separating  
line



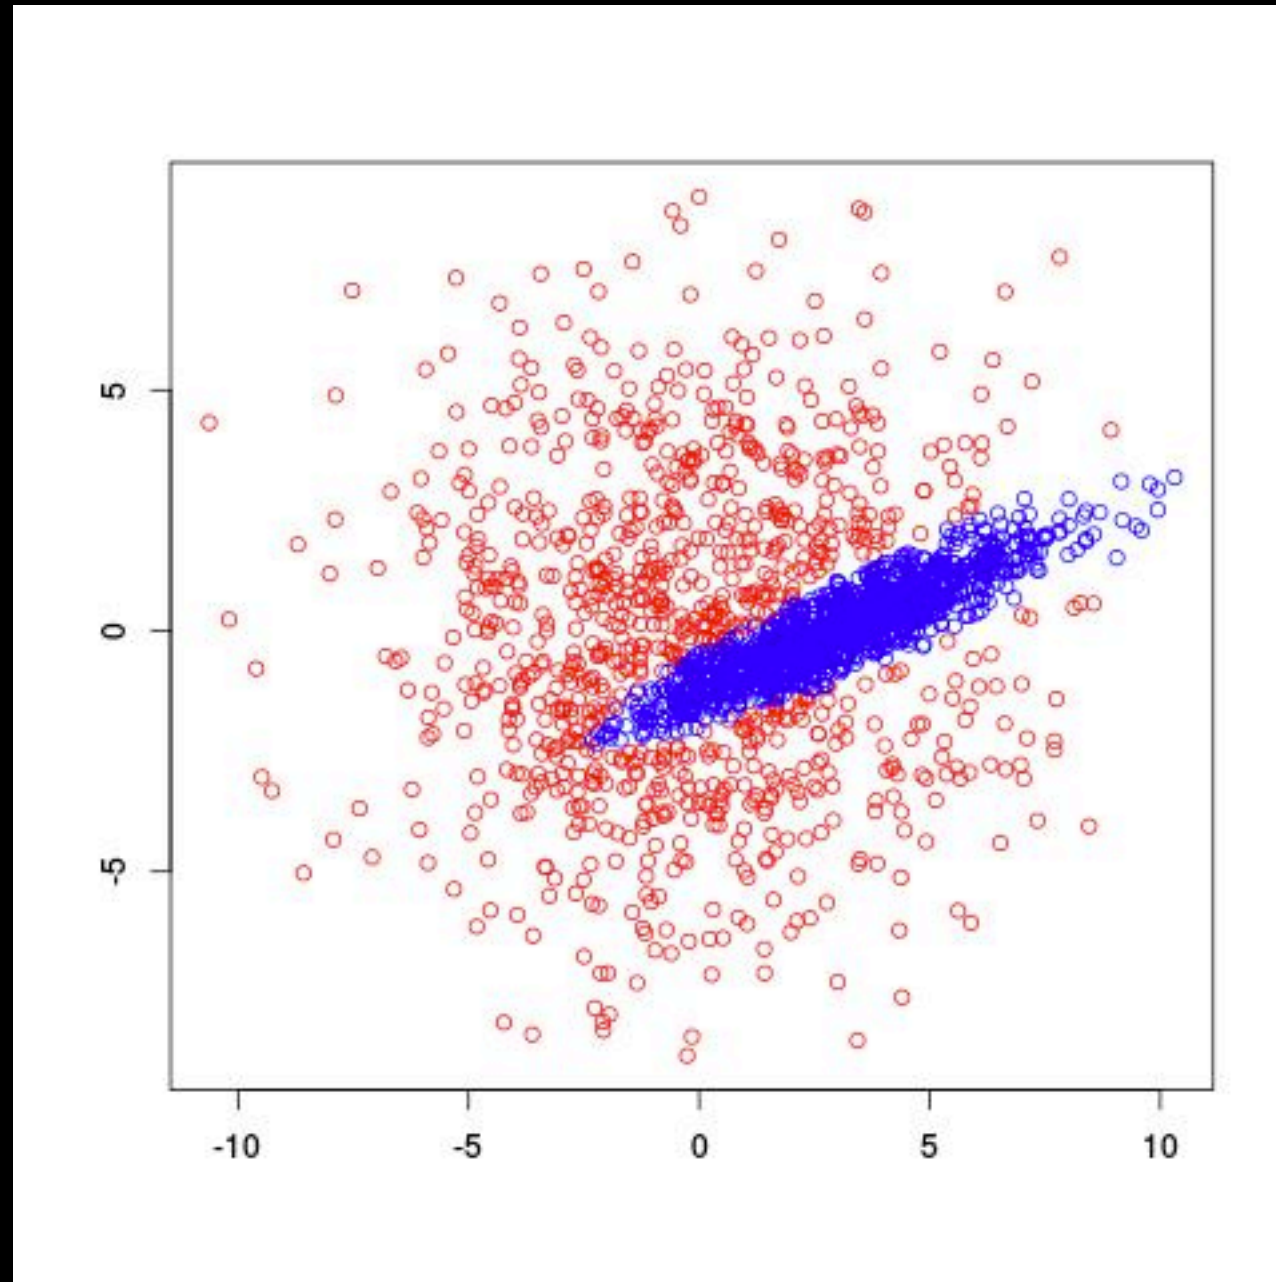
# K Nearest Neighbors - with Gaussian Mixture Models.



# K Nearest Neighbors - with Gaussian Mixture Models.



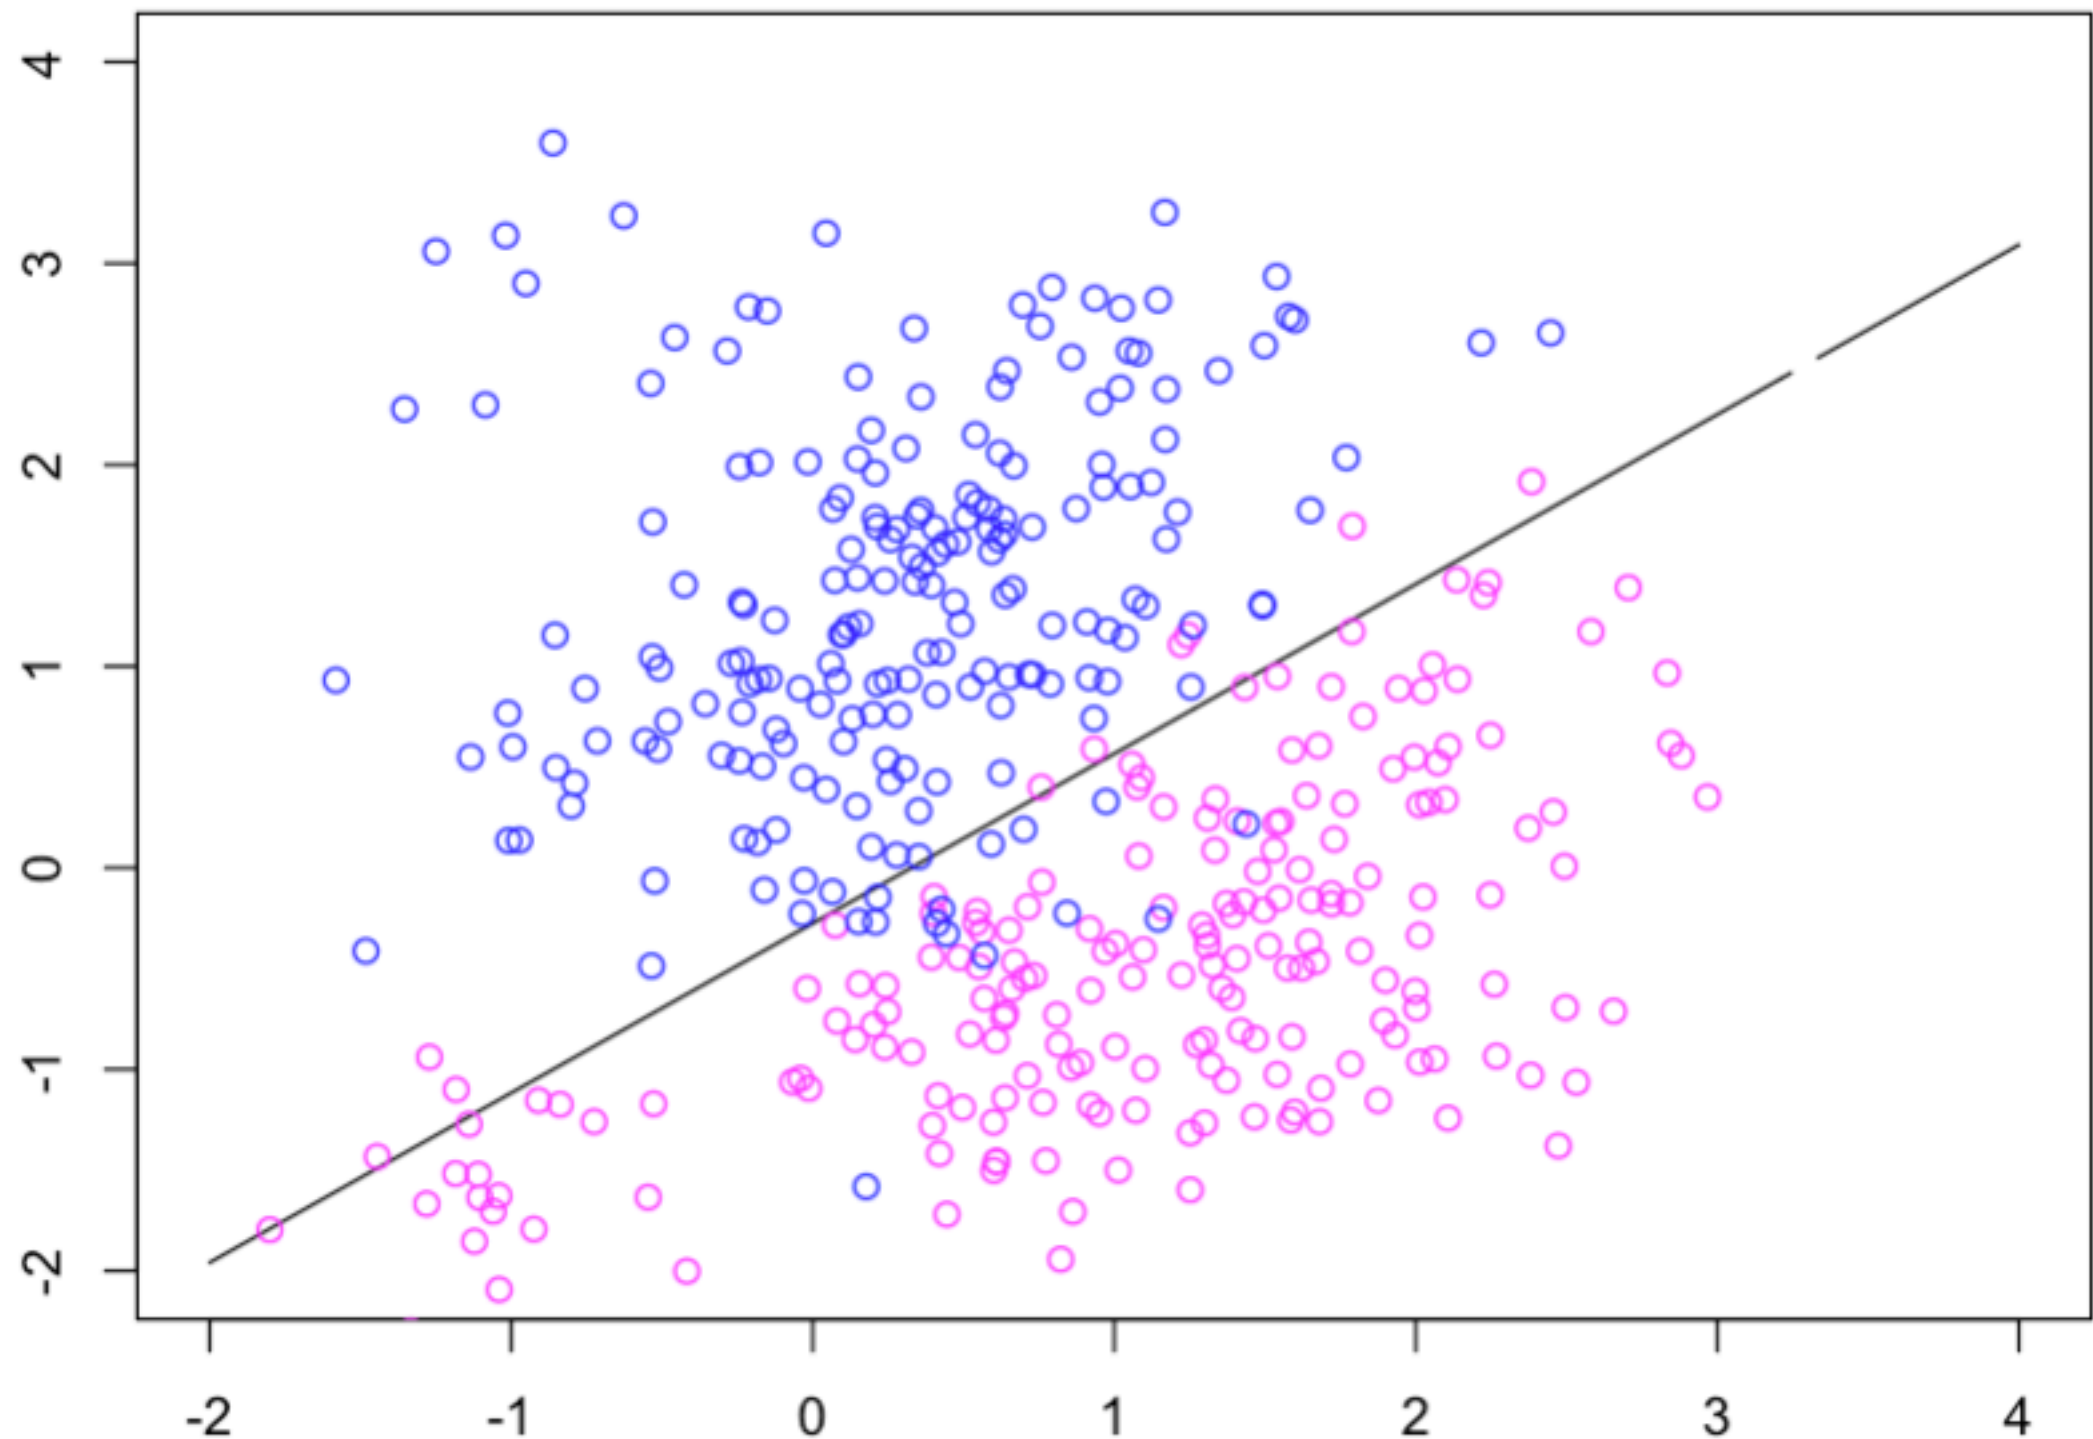
# K Nearest Neighbors - with Gaussian Mixture Models.



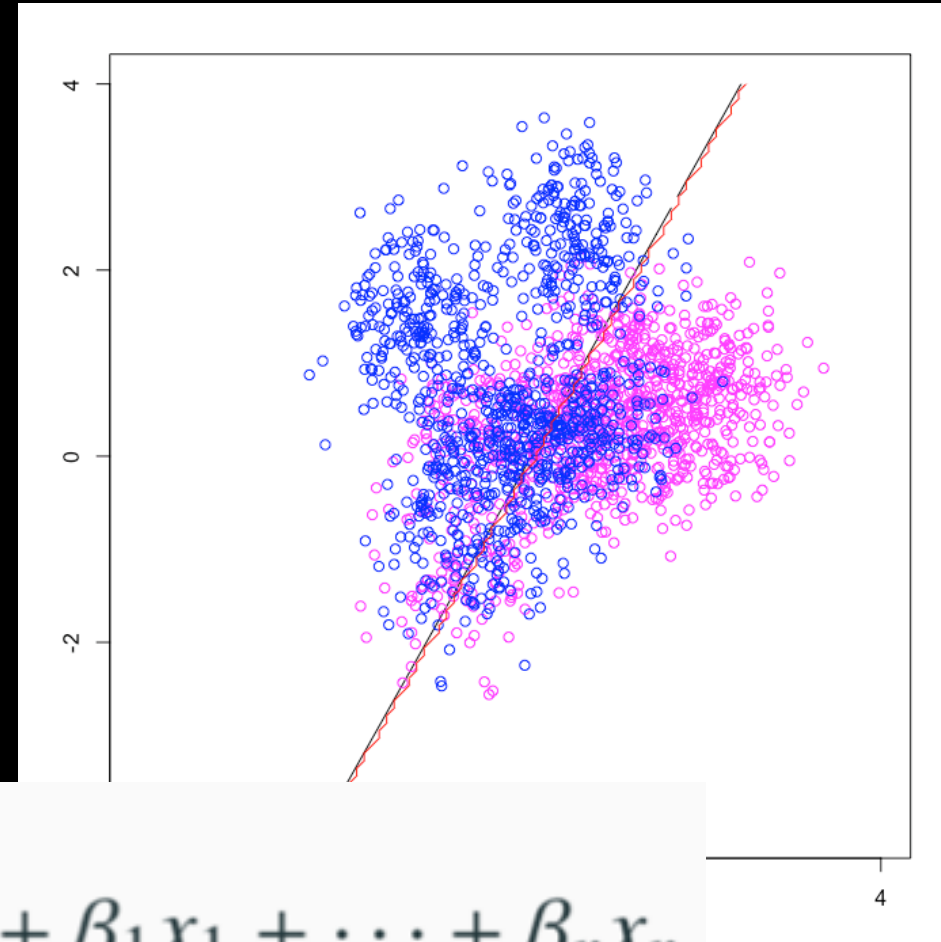
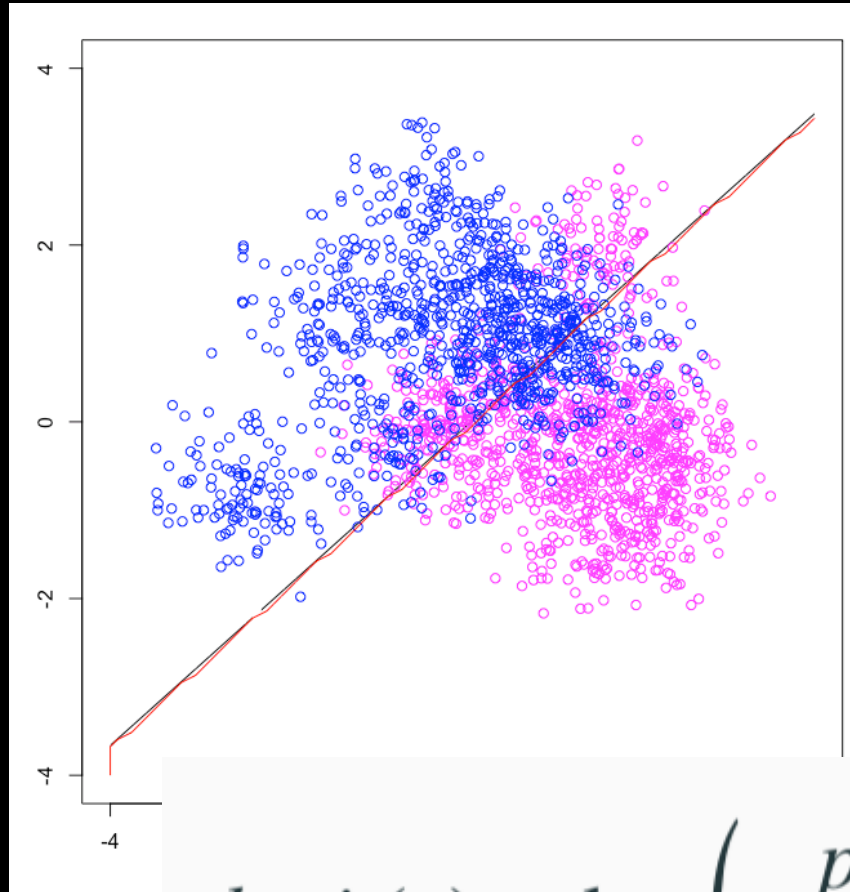
**K Nearest Neighbors, in R!**

# K Nearest Neighbors, in R!

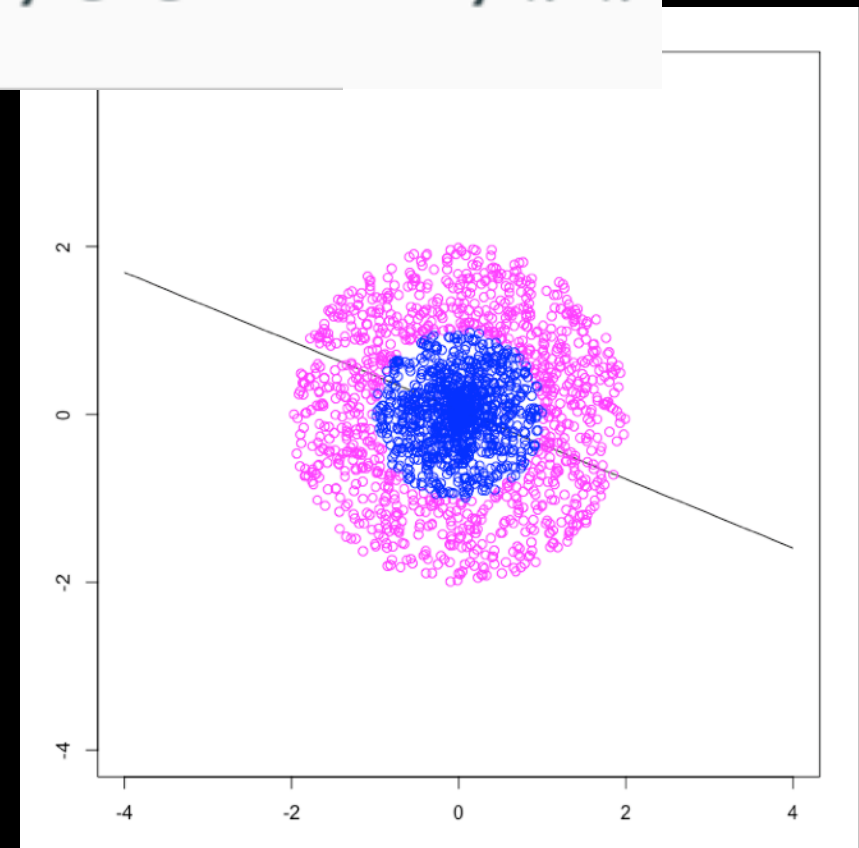
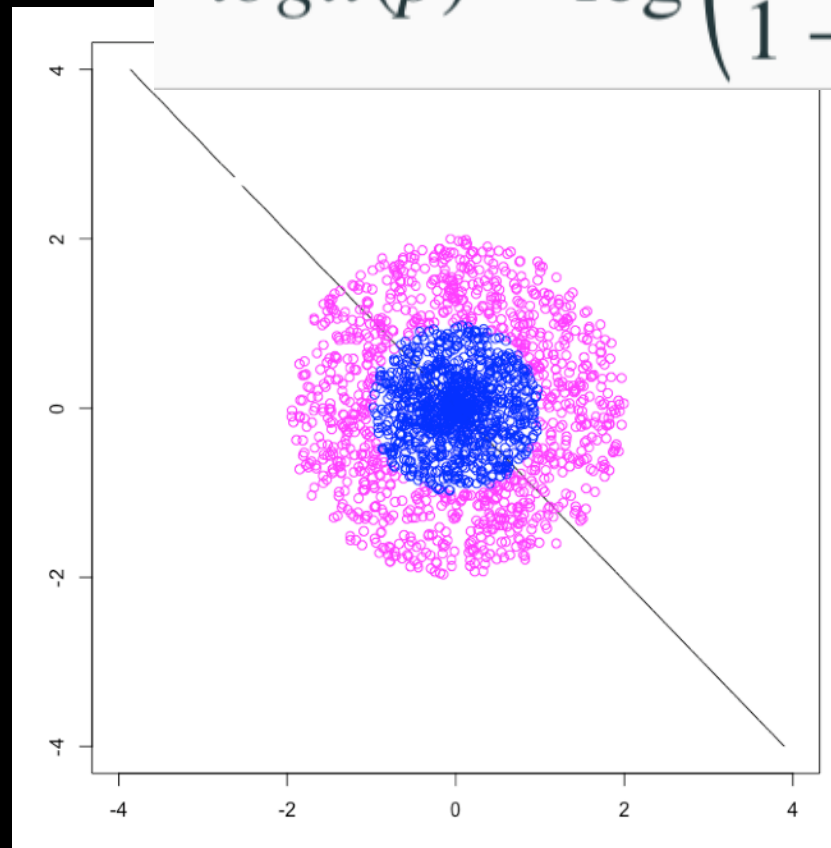
- just kidding logistic regression



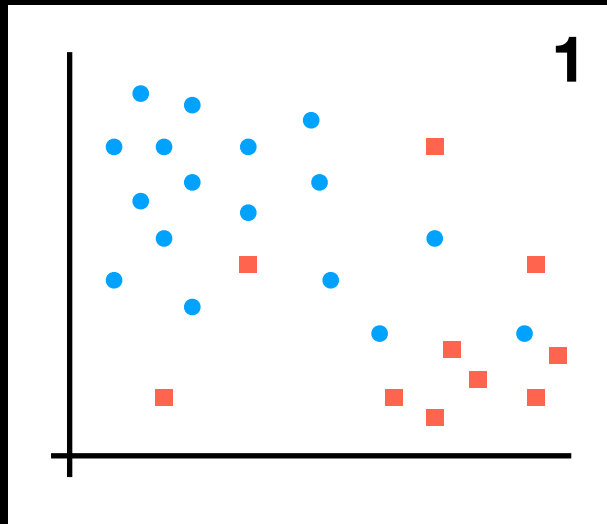
# GLM is clearly getting something wrong



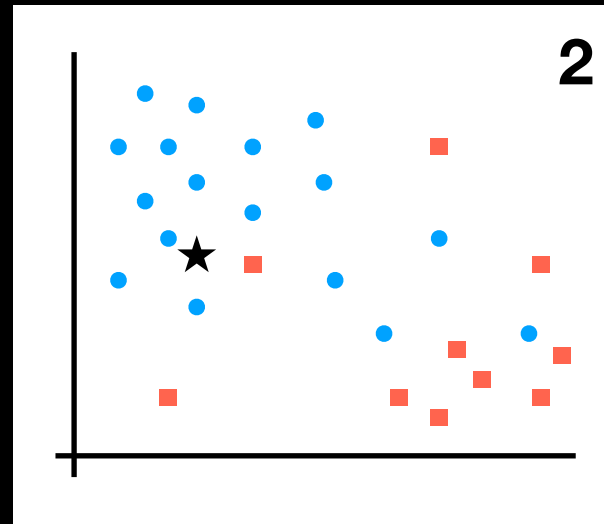
$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x_1 + \cdots + \beta_n x_n$$



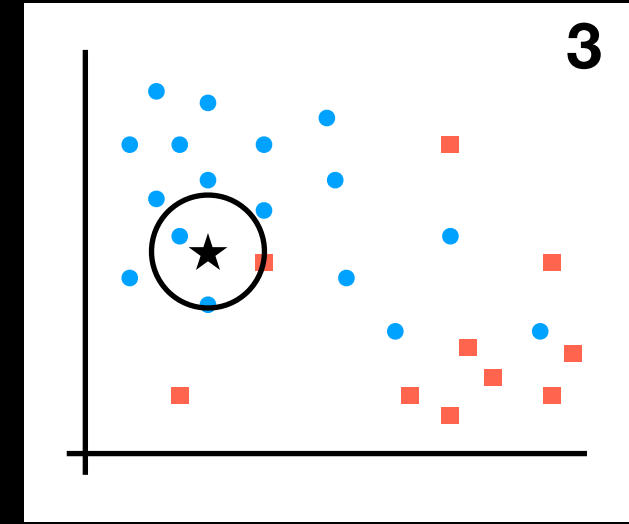
# K Nearest Neighbors, in pictures



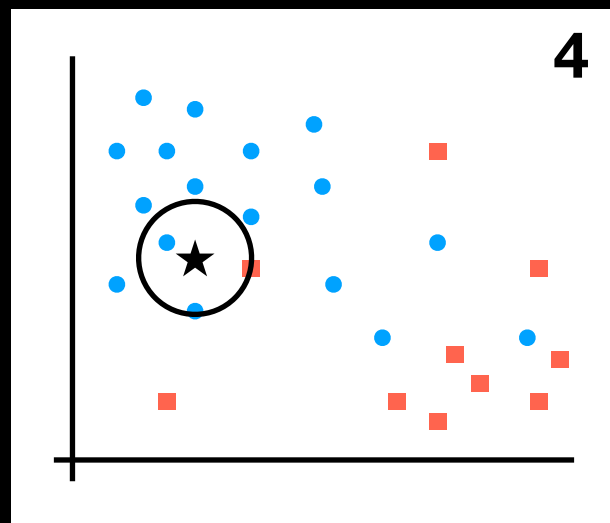
Sample (training) data  
representing underlying  
population



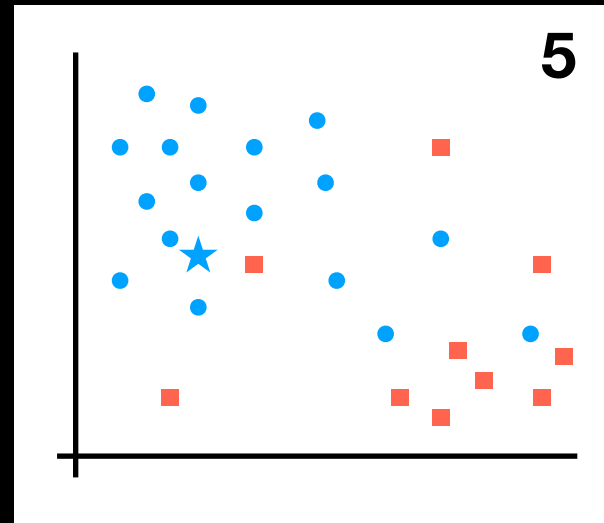
New point of interest



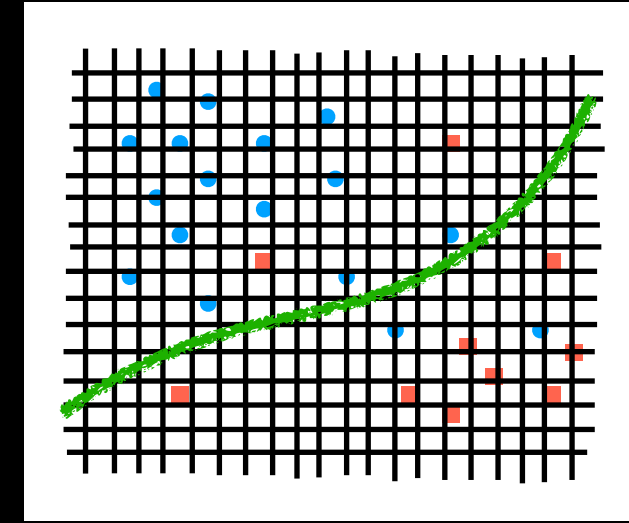
Find k nearest  
neighbors  
(here  $k = 3$ )



count "types" - here  
2/3 points are blue  
 $P(\text{blue}) = 2/3$   
 $P(\text{red}) = 1/3$



if  $P > \text{cut off}$  say new  
point is in that group  
here:  $P(\text{blue}) > 0.5$



Repeat 2-3 on a grid  
& draw a separating  
line

**K Nearest Neighbors, in R!**  
**For real this time!**



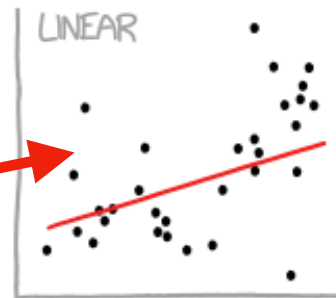
**Over/Under fitting - Quantifying how good your model is**

# CURVE-FITTING METHODS AND THE MESSAGES THEY SEND

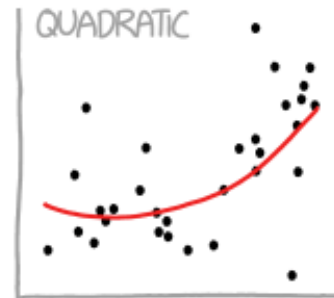
Overfit?

Overfit?

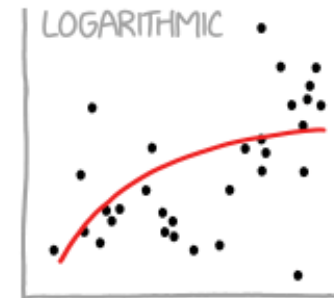
Overfit



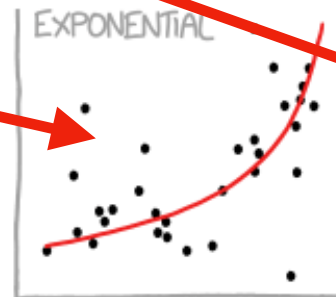
"HEY, I DID A REGRESSION."



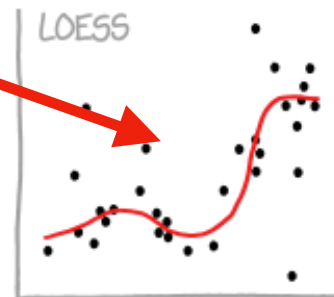
"I WANTED A CURVED LINE, SO I MADE ONE WITH MATH."



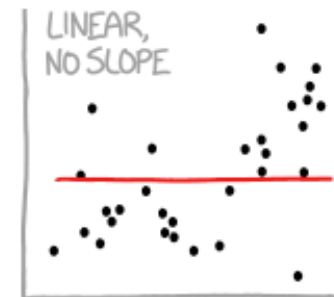
"LOOK, IT'S TAPERING OFF!"



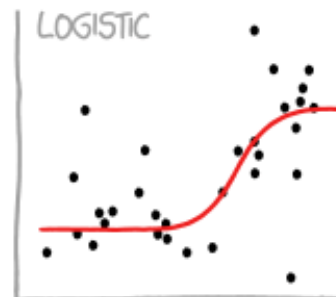
"LOOK, IT'S GROWING UNCONTROLLABLY!"



"I'M SOPHISTICATED, NOT LIKE THOSE BUMBLING POLYNOMIAL PEOPLE."



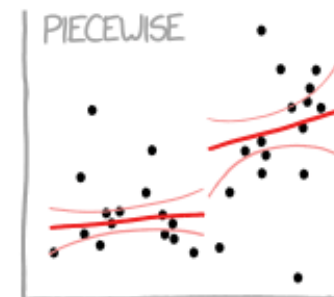
"I'M MAKING A SCATTER PLOT BUT I DON'T WANT TO."



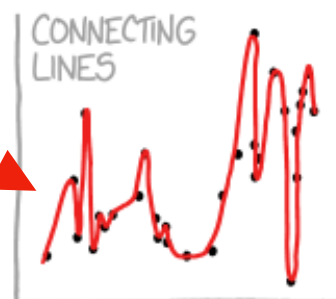
"I NEED TO CONNECT THESE TWO LINES, BUT MY FIRST IDEA DIDN'T HAVE ENOUGH MATH."



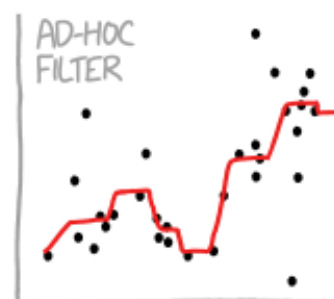
"LISTEN, SCIENCE IS HARD. BUT I'M A SERIOUS PERSON DOING MY BEST."



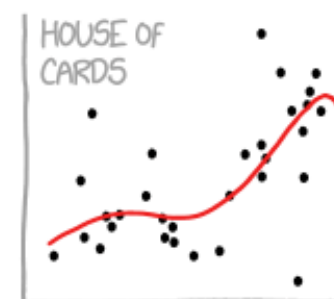
"I HAVE A THEORY, AND THIS IS THE ONLY DATA I COULD FIND."



"I CLICKED 'SMOOTH LINES' IN EXCEL."



"I HAD AN IDEA FOR HOW TO CLEAN UP THE DATA. WHAT DO YOU THINK?"



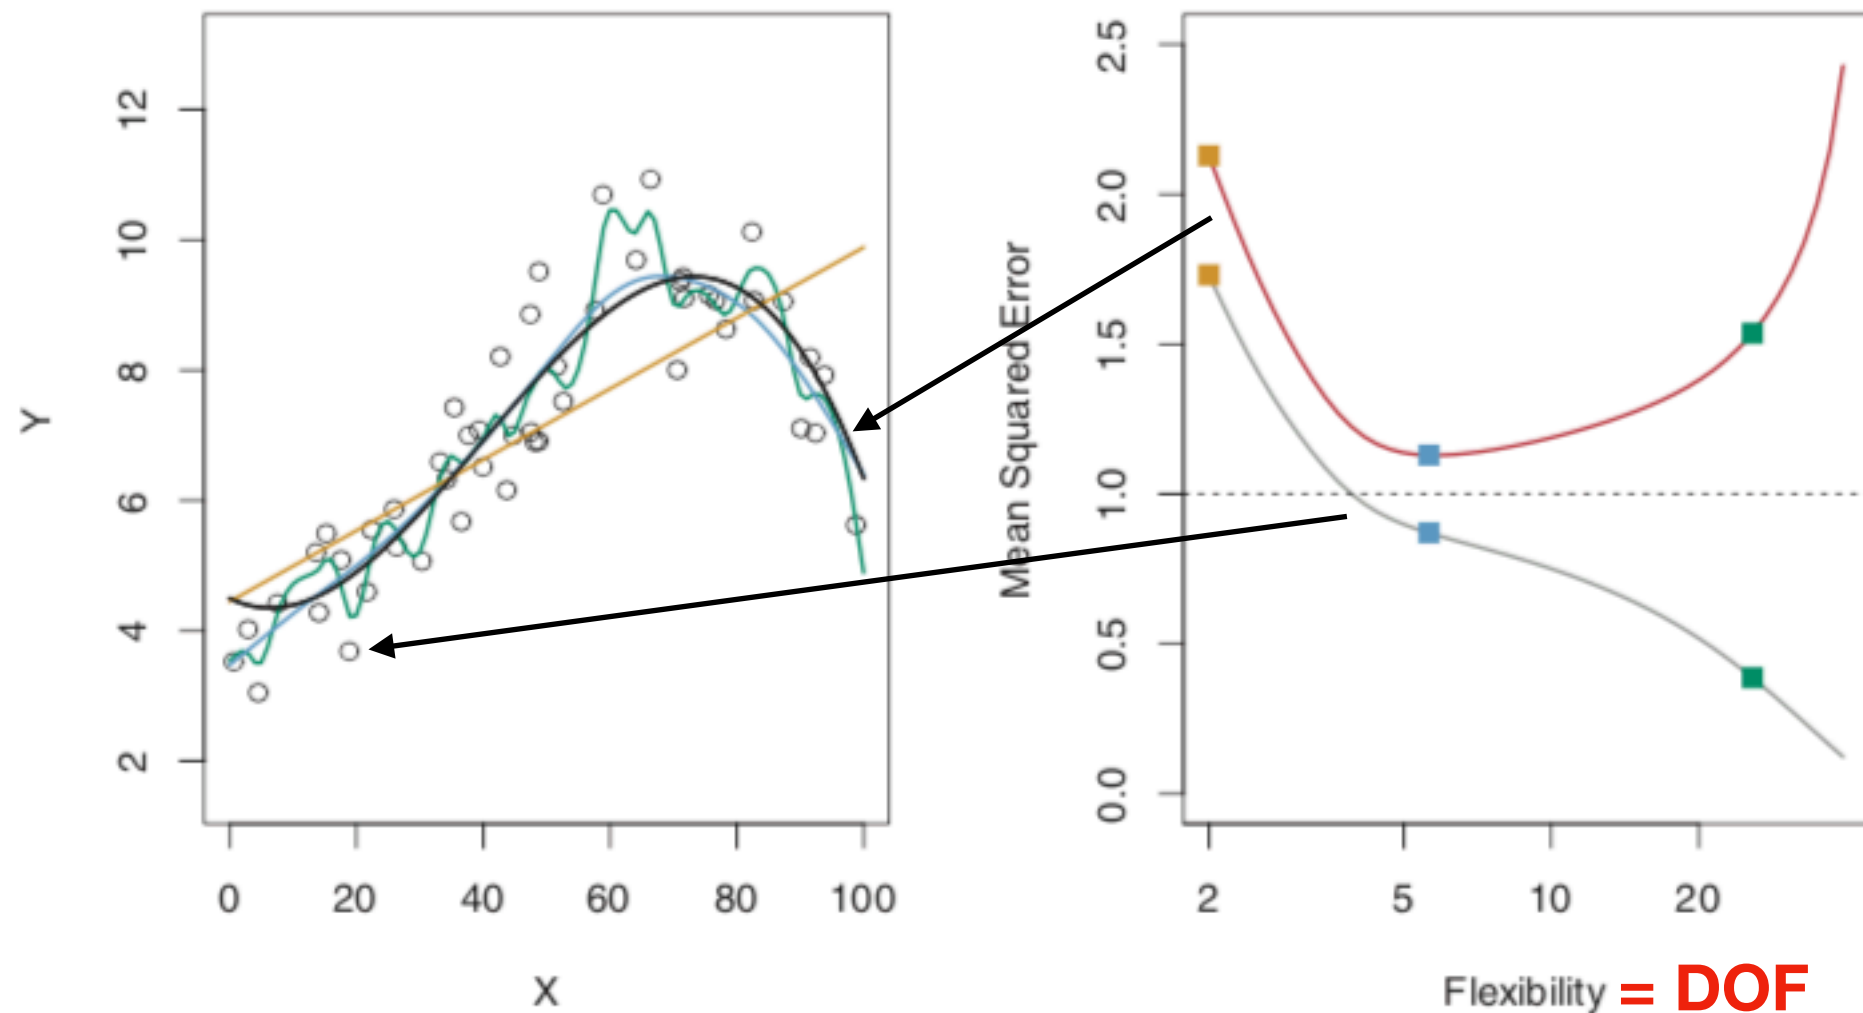
"AS YOU CAN SEE, THIS MODEL SMOOTHLY FITS THE— WAIT NO NO DON'T EXTEND IT AAAAAA!!!"

Underfit?

Is overfitting or underfitting worse?

# Bias-Variance Trade-Off (Second Glance)

$$E \left( y_0 - \hat{f}(x_0) \right)^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon).$$



- Actual underlying function -  $y$
- o Simulated data with added error ( $\epsilon$ )
- Linear fit
- Low “flexibility” smooth spline
- High “flexibility” smooth spline

# Bias-Variance Trade-Off (First Glance)

$$E \left( y_0 - \hat{f}(x_0) \right)^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon)$$

**mean square error** if we kept estimating response variable  $y$  by our fitted function of our explanatory variables,  $f(x)$  with different sample datasets at point  $x_0$

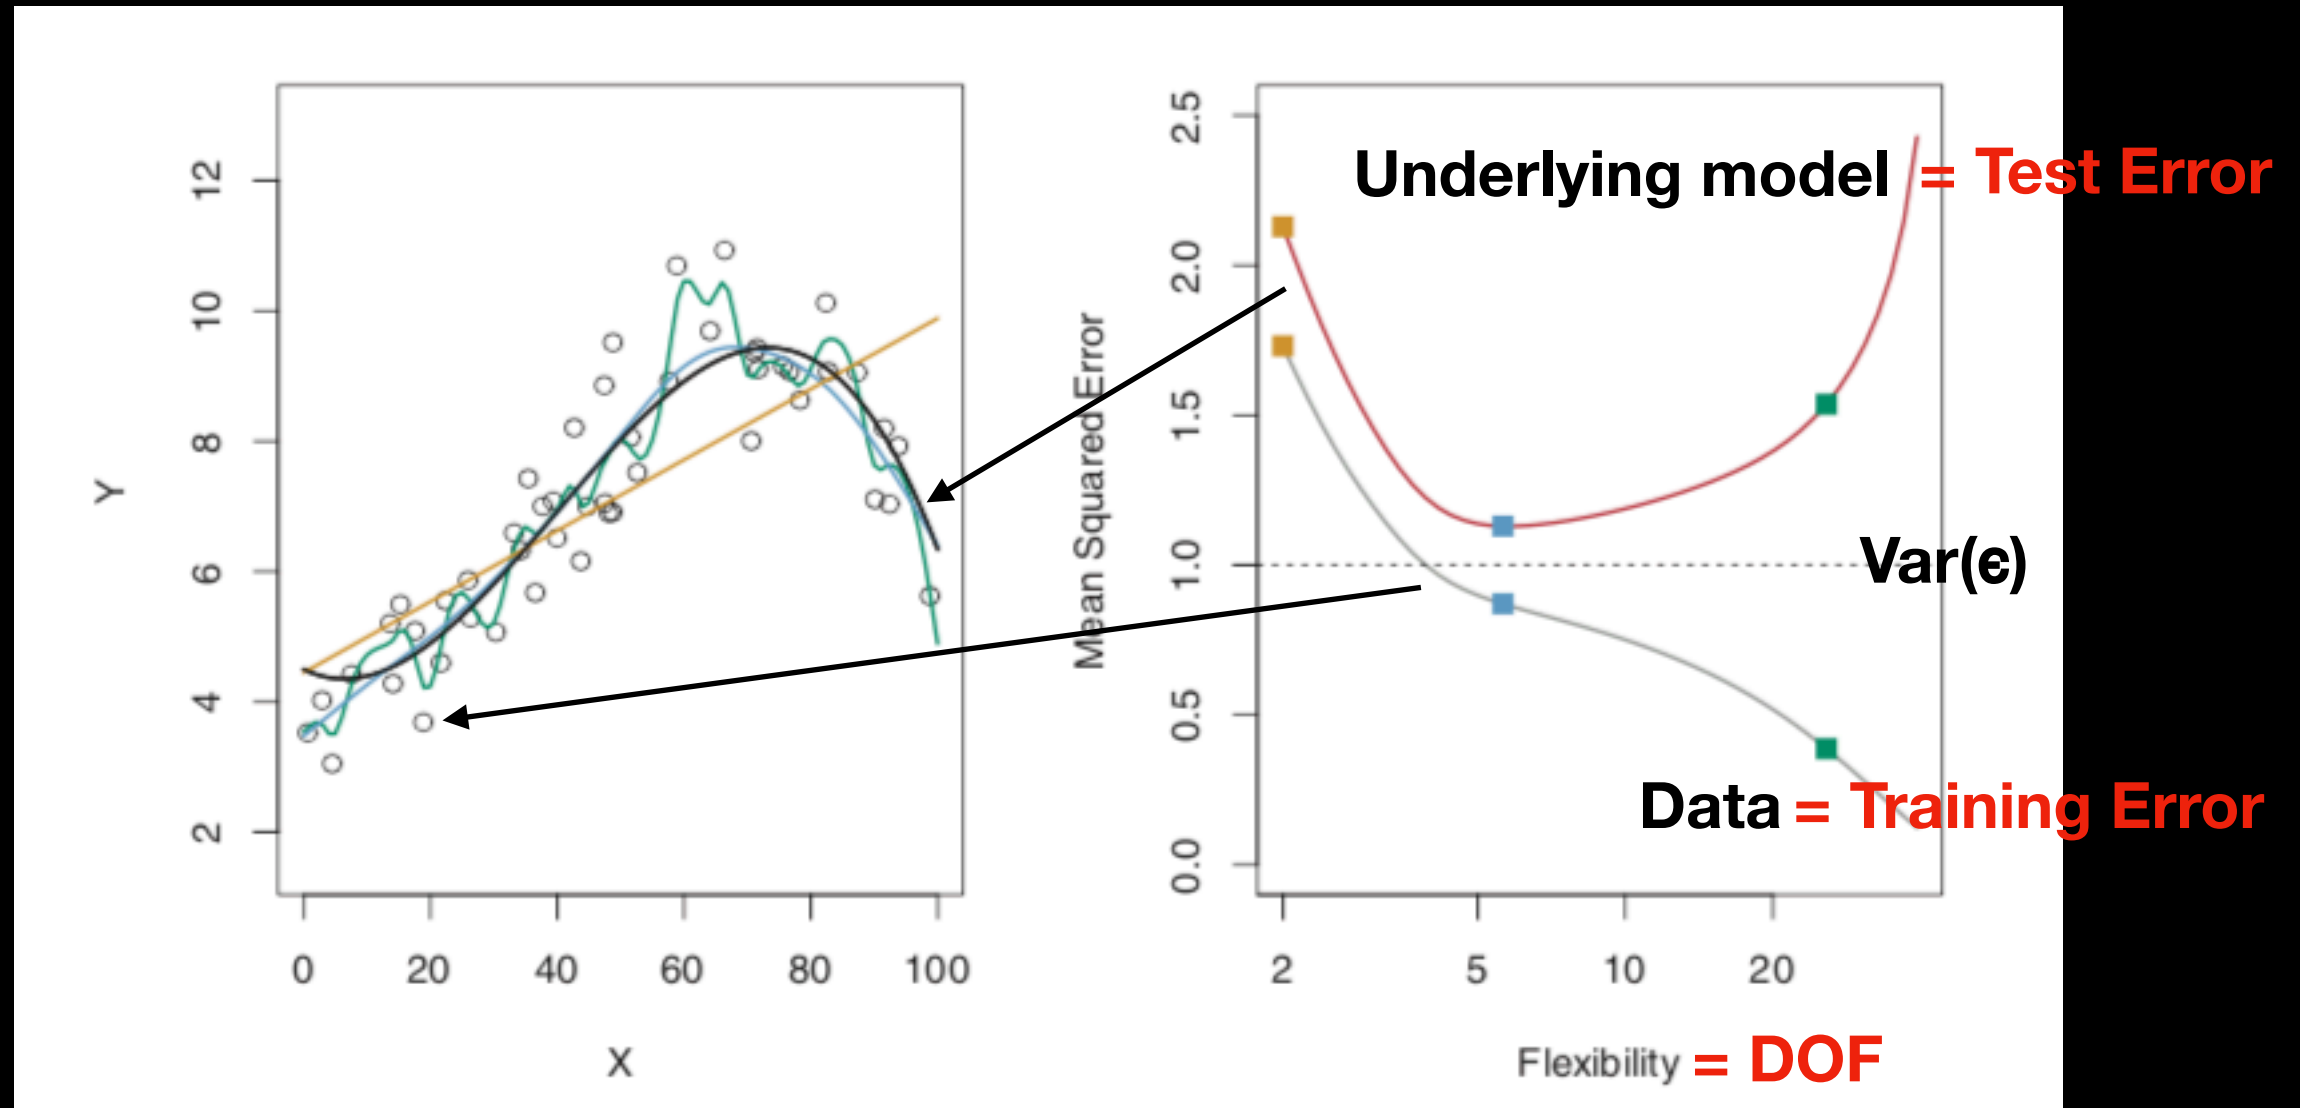
how much our function,  $f$ , changes if we use a different random sample  
(**variance**)

Inherent error (**bias**) in the fact that any model is only an approximation to reality

Inherent error in our measurements

# Bias-Variance Trade-Off (Second Glance)

$$E \left( y_0 - \hat{f}(x_0) \right)^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon).$$

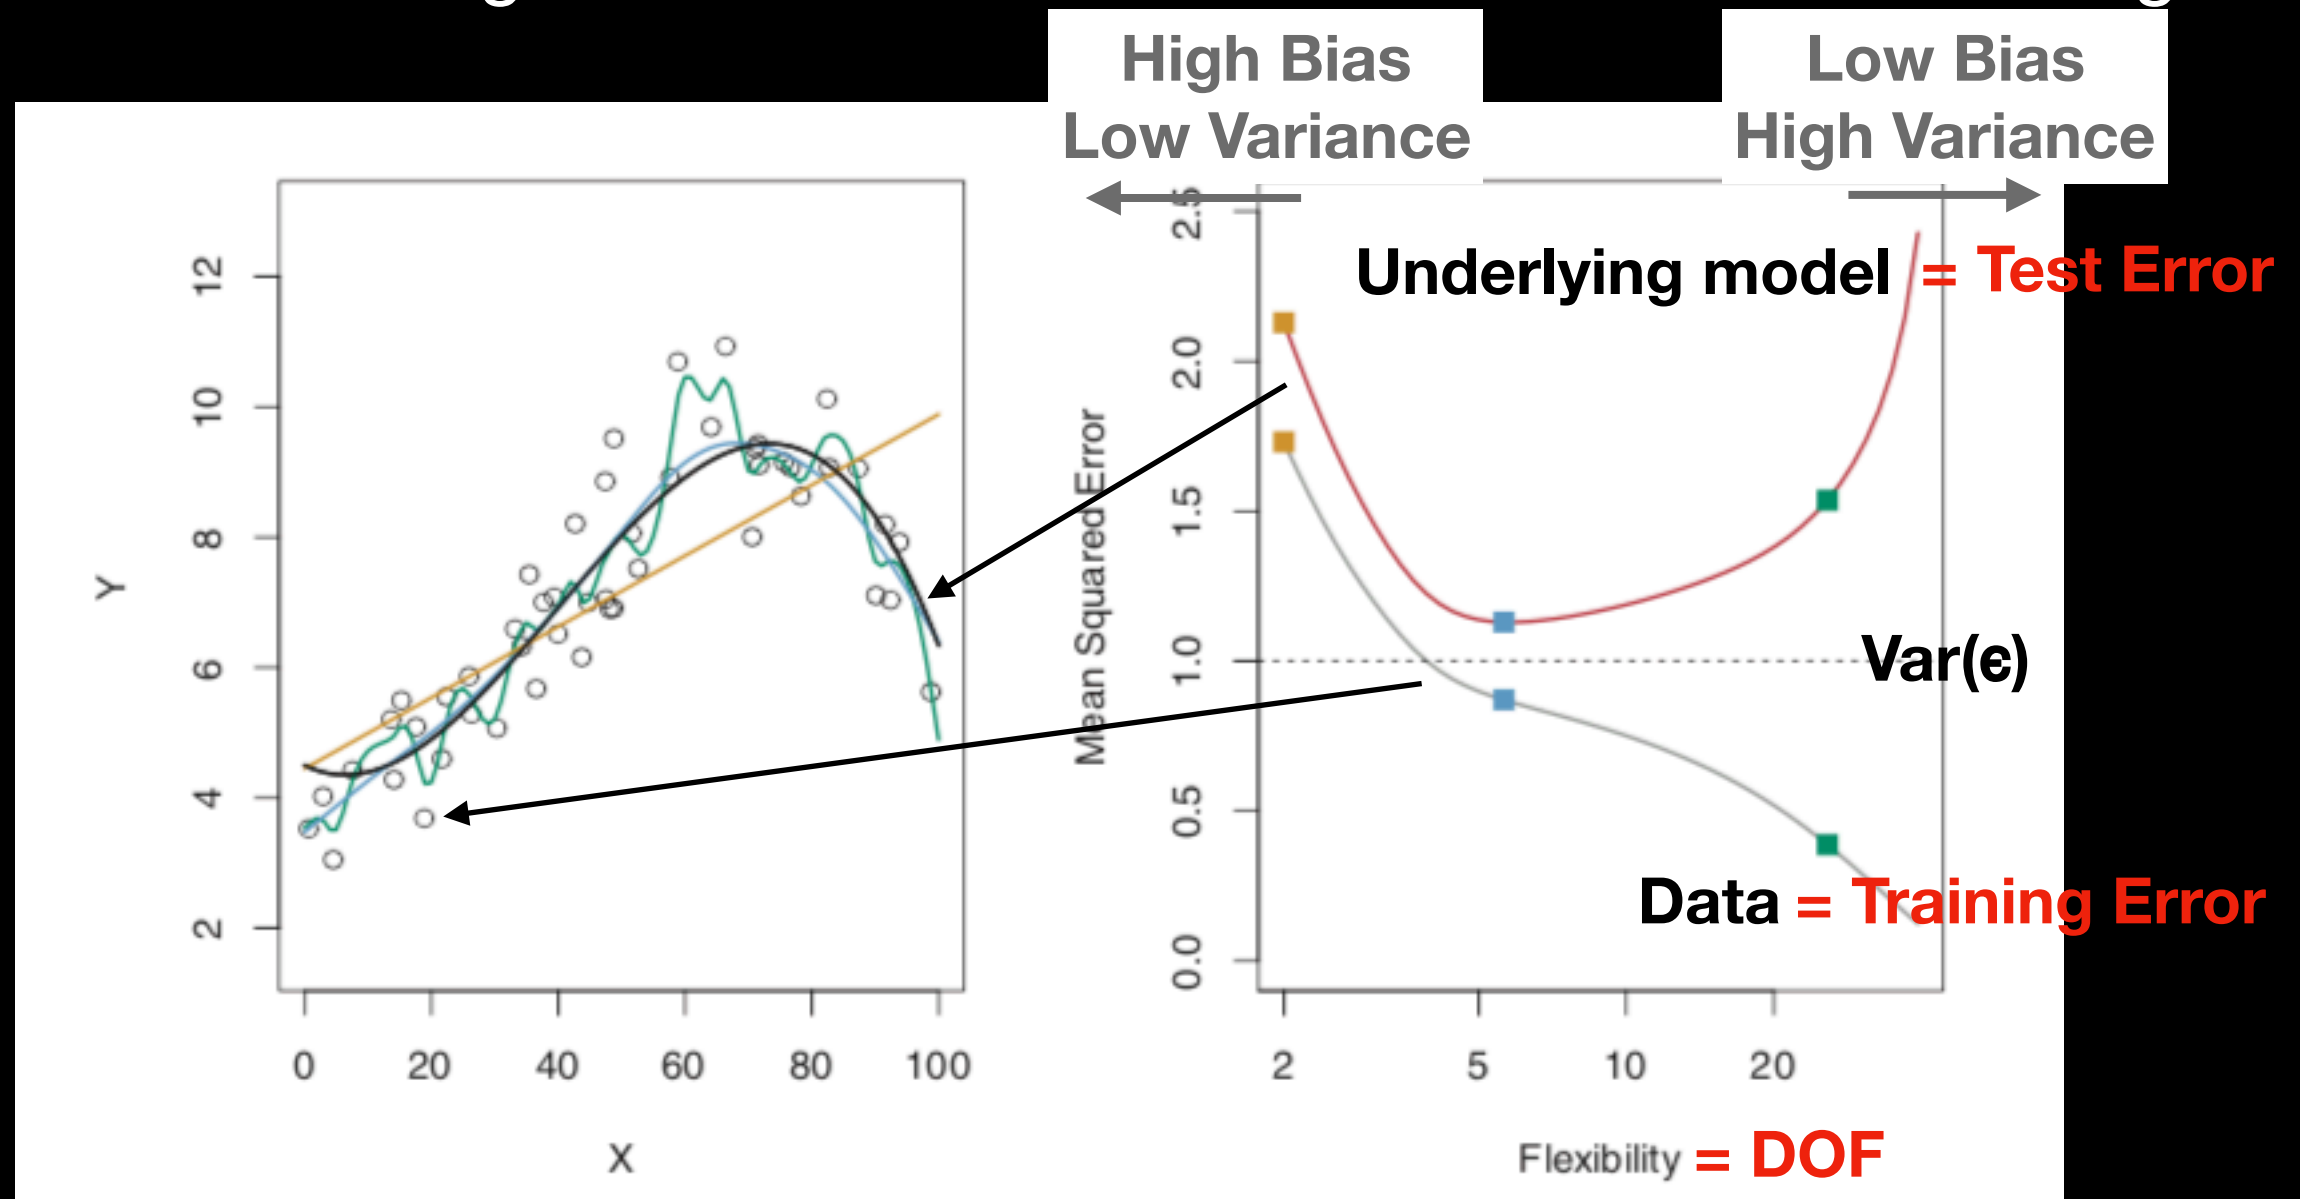


- Actual underlying function - y
- o Simulated data with added error ( $\epsilon$ )
- Linear fit
- Low “flexibility” smooth spline
- High “flexibility” smooth spline

# Bias-Variance Trade-Off (Second Glance)

$$E \left( y_0 - \hat{f}(x_0) \right)^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon).$$

- The **test error** is the average error that results from using a statistical learning method to predict the response on a new observation, one that was *not* used in training the method. (Hard to calculate w/o an underlying model.)
- In contrast, the **training error** can be easily calculated by applying the statistical learning method to the observations used in its training.



# Test & Training Error in KNN: With Math!

- The **test error** is the average error that results from using a statistical learning method to predict the response on a new observation, one that was *not* used in training the method. (Hard to calculate w/o an underlying model.)

$$\text{Ave} (I(y_0 \neq \hat{y}_0))$$

new observation, requires we know what that would be from an underlying model (or more observations)

What our calculated fit/model would predict

- In contrast, the **training error** can be easily calculated by applying the statistical learning method to the observations used in its training.

$$\frac{1}{n} \sum_{i=1}^n I(y_i \neq \hat{y}_i)$$

an observation in our current dataset

What our fit/model would predict

Here  $I = 1$  if  $y_i \neq \hat{y}_i$  and  $I = 0$  if  $y_i = \hat{y}_i$ , so larger  $I$  means worse model

# K Nearest Neighbors (and MLR) → Cross-Validation & Bootstrap

p-values quantify the fit of  
individual parameters

quantify how good the  
*model* is

But first: some definitions!

Using our KNN example in R with an underlying model!